

Program:

### #Import Library

```
import pandas as pd
```

```
import numpy as np
```

```
import warnings
```

```
warnings.simplefilter("ignore")
```

### # Load the dataset

```
df=pd.read_csv("D:\Sethu 45\Class prediction 04\Traffic.csv")
```

```
df.head()
```

### Output:

	Time	Date	Day of the week	CarCount	BikeCount	BusCount	TruckCount	Total	Traffic Situation
0	12:00:00 AM	10	Tuesday	31	0	4	4	39	low
1	12:15:00 AM	10	Tuesday	49	0	3	3	55	low
2	12:30:00 AM	10	Tuesday	46	0	3	6	55	low
3	12:45:00 AM	10	Tuesday	51	0	2	5	58	low
4	1:00:00 AM	10	Tuesday	57	6	15	16	94	normal

```
df.info()
```

### Output:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2976 entries, 0 to 2975
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Time                  2976 non-null   object
1   Date                  2976 non-null   int64
2   Day of the week       2976 non-null   object
3   CarCount              2976 non-null   int64
4   BikeCount             2976 non-null   int64
5   BusCount              2976 non-null   int64
6   TruckCount            2976 non-null   int64
7   Total                 2976 non-null   int64
8   Traffic Situation     2976 non-null   object
dtypes: int64(6), object(3)
memory usage: 209.4+ KB
```

```
df.describe()
```

**Output:**

	Date	CarCount	BikeCount	BusCount	TruckCount	Total
count	2976.000000	2976.000000	2976.000000	2976.000000	2976.000000	2976.000000
mean	16.000000	68.696573	14.917339	15.279570	15.324933	114.218414
std	8.945775	45.850693	12.847518	14.341986	10.603833	60.190627
min	1.000000	6.000000	0.000000	0.000000	0.000000	21.000000
25%	8.000000	19.000000	5.000000	1.000000	6.000000	55.000000
50%	16.000000	64.000000	12.000000	12.000000	14.000000	109.000000
75%	24.000000	107.000000	22.000000	25.000000	23.000000	164.000000
max	31.000000	180.000000	70.000000	50.000000	40.000000	279.000000

**# Explore and preprocess the dataset**

```
df.isnull().sum()
```

**Output:**

```
Time          0
Date          0
Day of the week  0
CarCount      0
BikeCount     0
BusCount      0
TruckCount    0
Total         0
Traffic Situation 0
dtype: int64
```

```
df=df.drop('Time',axis=1)
```

```
df.head()
```

**Output:**

	Date	Day of the week	CarCount	BikeCount	BusCount	TruckCount	Total	Traffic Situation
0	10	Tuesday	31	0	4	4	39	low
1	10	Tuesday	49	0	3	3	55	low
2	10	Tuesday	46	0	3	6	55	low
3	10	Tuesday	51	0	2	5	58	low
4	10	Tuesday	57	6	15	16	94	normal

### # Encoding the categorical values

```
from sklearn.preprocessing import LabelEncoder  
label_encoder = LabelEncoder()  
categorical_columns = ['Day of the week', 'Traffic Situation']  
df[categorical_columns] = df[categorical_columns].apply(label_encoder.fit_transform)  
df[categorical_columns]
```

### Output:

	Day of the week	Traffic Situation
0	5	2
1	5	2
2	5	2
3	5	2
4	5	3
...	...	...
2971	4	3
2972	4	3
2973	4	3
2974	4	3
2975	4	3

2976 rows x 2 columns

### #Split the dataset into features(x)

```
x=df.iloc[:, :-1]
```

x

## Output:

	Date	Day of the week	CarCount	BikeCount	BusCount	TruckCount	Total
0	10	5	31	0	4	4	39
1	10	5	49	0	3	3	55
2	10	5	46	0	3	6	55
3	10	5	51	0	2	5	58
4	10	5	57	6	15	16	94
...	...	...	...	...	...	...	...
2971	9	4	16	3	1	36	56
2972	9	4	11	0	1	30	42
2973	9	4	15	4	1	25	45
2974	9	4	16	5	0	27	48
2975	9	4	14	3	1	15	33

2976 rows x 7 columns

## #Split the variables into target variables(y)

```
y=df.iloc[:,-1]
```

y

## Output:

```
0      2
1      2
2      2
3      2
4      3
..
2971    3
2972    3
2973    3
2974    3
2975    3
Name: Traffic Situation, Length: 2976, dtype: int32
```

## # Further split the dataset into training and testing sets

```
from sklearn. model_selection import train_test_split
```

```
x_train, x_test, y_train, y_test=train_test_split (x, y, random_state=0)
```

### # Feature scaling the data

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler ()
```

```
x_train = scaler.fit_transform(x_train)
```

```
x_test = scaler.transform(x_test)
```

```
x_train
```

### Output:

```
array([[ -0.33600687, -1.59798223,  1.98010887, ...,  1.00899718,
        -1.1633432 ,  2.0733575 ],
       [  0.43635908, -1.59798223,  2.30697357, ..., -0.31203497,
        -1.1633432 ,  2.20601196],
       [  1.42940101, -0.10771914,  1.10846967, ...,  0.66135714,
        -0.68866549,  0.8628855 ],
       ...,
       [  1.20872503, -1.59798223, -0.351526 , ..., -0.590147 ,
        -0.02411669, -0.48024096],
       [-1.10837282, -1.10122786, -1.15779226, ..., -1.07684306,
         0.83030319, -1.1766769 ],
       [-0.99803483,  0.88578959,  0.01892066, ..., -0.72920302,
         1.30498091, -0.09885937]])
```

```
x_test
```

### Output:

```
array([[ -0.99803483,  0.88578959, -1.24495618, ..., -1.07684306,
         1.21004537, -1.19325871],
       [-1.66006278,  1.38254396,  0.38936732, ...,  0.17466108,
         1.39991645,  0.43175849],
       [-1.32904881, -0.6044735 , -0.22078012, ...,  1.2175812 ,
         0.92523874,  0.21619498],
       ...,
       [  1.31906302, -0.6044735 , -1.07062834, ..., -0.17297896,
         0.64043211, -0.91136797],
       [  0.32602109,  0.38903523, -1.13600128, ..., -1.00731505,
         0.45056102, -1.20984052],
       [-1.4393868 , -1.59798223, -0.351526 , ..., -0.79873103,
        -0.02411669, -0.57973181]])
```

### #Initialise the Logistic Regression Model

```
from sklearn.linear_model import LogisticRegression
```

```
model = LogisticRegression ()
```

## #Train the Logistic Regression Model

```
model.fit(x_train,y_train)
```

**Output:**

```
LogisticRegression
```

## #Make predictions

```
y_pred=model.predict(x_test)
```

y\_pred

**Output:**

```
array([3, 1, 1, 0, 3, 3, 3, 0, 0, 3, 0, 3, 0, 3, 3, 0, 1, 0, 2, 3, 3, 1,
       3, 2, 3, 0, 3, 3, 3, 0, 0, 3, 3, 3, 0, 0, 3, 0, 1, 3, 0, 0, 2, 3,
       0, 3, 3, 0, 3, 3, 3, 0, 3, 0, 3, 0, 3, 0, 3, 3, 0, 3, 3, 2, 0,
       3, 3, 2, 0, 1, 0, 2, 3, 3, 2, 3, 0, 3, 3, 3, 0, 0, 3, 3, 0, 2, 3,
       3, 0, 0, 3, 2, 3, 3, 0, 3, 1, 1, 3, 0, 0, 3, 3, 3, 0, 3, 3, 3, 3,
       3, 0, 3, 1, 3, 3, 0, 3, 3, 0, 3, 3, 3, 2, 3, 0, 3, 0, 3, 0, 3,
       0, 3, 3, 3, 3, 0, 3, 3, 0, 3, 1, 3, 2, 2, 3, 0, 3, 3, 0, 3, 3, 2,
       3, 2, 2, 0, 3, 3, 2, 3, 0, 3, 3, 3, 3, 3, 1, 3, 3, 2, 3, 2, 0, 3,
       0, 3, 3, 0, 3, 3, 0, 3, 0, 1, 1, 3, 0, 0, 3, 3, 2, 0, 0, 0, 2, 3,
       0, 3, 3, 0, 1, 1, 3, 3, 3, 1, 3, 3, 3, 1, 0, 0, 0, 2, 2, 1, 3, 0,
       3, 3, 3, 3, 0, 3, 1, 3, 3, 0, 3, 3, 3, 0, 3, 3, 2, 3, 0, 3, 3, 3,
       3, 3, 3, 1, 1, 1, 1, 3, 0, 3, 3, 3, 3, 0, 3, 3, 0, 3, 0, 3, 3, 0,
       3, 3, 3, 1, 0, 1, 2, 3, 3, 2, 0, 3, 3, 3, 0, 1, 0, 3, 3, 3, 1, 3,
       0, 3, 1, 3, 0, 3, 3, 3, 3, 3, 0, 0, 2, 0, 0, 3, 2, 2, 3, 0, 3,
       3, 0, 3, 3, 3, 3, 1, 0, 3, 0, 0, 2, 1, 2, 0, 3, 3, 3, 0, 3, 3, 0,
       3, 1, 3, 3, 0, 0, 0, 0, 3, 3, 3, 3, 0, 3, 3, 3, 3, 1, 3, 3, 0,
       3, 3, 3, 2, 3, 3, 3, 0, 0, 3, 0, 3, 3, 3, 3, 3, 3, 0, 3, 0, 3, 3,
       0, 3, 3, 0, 3, 3, 1, 0, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 0, 3,
       0, 3, 3, 2, 3, 2, 0, 3, 3, 3, 0, 0, 3, 3, 3, 1, 3, 3, 3, 3, 2, 3,
       3, 3, 3, 3, 2, 2, 2, 3, 0, 3, 3, 1, 1, 0, 2, 3, 3, 2, 3, 3, 0, 3,
       3, 3, 3, 3, 0, 1, 3, 3, 2, 3, 0, 3, 1, 3, 0, 3, 3, 3, 2, 3, 3,
       3, 2, 2, 3, 3, 3, 1, 0, 3, 3, 2, 3, 0, 3, 0, 3, 2, 3, 3, 3, 3, 0,
       2, 3, 3, 3, 2, 3, 3, 3, 3, 3, 0, 0, 3, 0, 1, 0, 3, 3, 1, 3, 3,
       0, 3, 3, 2, 3, 3, 1, 3, 0, 3, 1, 2, 3, 3, 3, 3, 0, 3, 3, 3, 1, 3,
       0, 3, 3, 3, 3, 0, 2, 2, 0, 2, 0, 0, 3, 0, 1, 0, 3, 3, 0, 0, 3,
       3, 3, 3, 0, 3, 3, 3, 1, 0, 0, 0, 3, 1, 3, 0, 3, 3, 3, 3, 3, 3,
       3, 2, 3, 0, 3, 3, 1, 3, 3, 3, 0, 0, 0, 3, 3, 3, 2, 2, 3, 3, 3, 3,
       1, 3, 2, 3, 3, 3, 3, 3, 0, 0, 3, 2, 1, 0, 1, 3, 3, 0, 3, 3, 3,
       1, 0, 3, 3, 3, 2, 0, 3, 0, 2, 3, 0, 3, 3, 3, 3, 0, 3, 1, 3, 3, 0,
       0, 3, 3, 0, 3, 3, 3, 3, 3, 0, 0, 3, 3, 3, 3, 3, 3, 3, 3, 0, 3,
       2, 3, 1, 3, 3, 3, 0, 1, 0, 0, 3, 3, 3, 0, 3, 3, 3, 1, 2, 3, 3, 0,
       3, 3, 3, 3, 3, 0, 2, 0, 3, 3, 0, 0, 3, 3, 3, 3, 2, 3, 3, 3, 0, 1,
       3, 2, 1, 3, 3, 2, 3, 3, 1, 3, 0, 3, 3, 1, 3, 3, 3, 1, 2, 3, 0,
       0, 3, 0, 3, 3, 2, 3, 3, 0, 3, 3, 3, 0, 3, 0, 3, 3, 3])
```

## #Evaluate the model

```
from sklearn.metrics import accuracy_score, classification_report
```

```
accuracy = accuracy_score(y_test, y_pred) * 100
```

```
print ("Accuracy of the model is {:.2f}".format(accuracy))
```

### Output:

```
Accuracy of the model is 88.44
```

### # Classification Report

```
from sklearn.metrics import classification_report
class_report = classification_report(y_test, y_pred)
print(f'\nClassification Report:\n{class_report}')
```

### Output:

```
Classification Report:
              precision    recall  f1-score   support

     0           0.92       0.98       0.95         168
     1           0.85       0.65       0.74          80
     2           0.73       0.67       0.70          73
     3           0.90       0.93       0.91         423

 accuracy          0.88         0.88         0.88         744
 macro avg          0.85         0.81         0.82         744
 weighted avg       0.88         0.88         0.88         744
```

### #Predict for new data

```
new_data = [9, 5, 6, 3, 2, 5, 45]
predictions = model.predict([new_data])
print(predictions)
```

### Output:

```
[0]
```

### # Decoded the output of prediction

```
decoded_predictions = label_encoder.inverse_transform(predictions)
print(decoded_predictions)
```

### Output:

```
['heavy']
```



Program:

### #Import Library

```
import pandas as pd
```

```
import numpy as np
```

```
import warnings
```

```
warnings.simplefilter("ignore")
```

### # Load the dataset

```
df=pd.read_csv("D:\Sethu 45\Class prediction 04\Trip.csv")
```

```
df.head()
```

### Output:

	Duration	Start date	End date	Start station number	Start station	End station number	End station	Bike number	Member type
0	1012	20-09-2010 11:27	20-09-2010 11:43	31208	M St & New Jersey Ave SE	31108	4th & M St SW	W00742	Member
1	61	20-09-2010 11:41	20-09-2010 11:42	31209	1st & N St SE	31209	1st & N St SE	W00032	Member
2	2690	20-09-2010 12:05	20-09-2010 12:50	31600	5th & K St NW	31100	19th St & Pennsylvania Ave NW	W00993	Member
3	1406	20-09-2010 12:06	20-09-2010 12:29	31600	5th & K St NW	31602	Park Rd & Holmead Pl NW	W00344	Member
4	1413	20-09-2010 12:10	20-09-2010 12:34	31100	19th St & Pennsylvania Ave NW	31201	15th & P St NW	W00883	Member

```
df.info()
```

### Output:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 115597 entries, 0 to 115596
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Duration              115597 non-null  int64
1   Start date            115597 non-null  object
2   End date              115597 non-null  object
3   Start station number  115597 non-null  int64
4   Start station         115597 non-null  object
5   End station number    115597 non-null  int64
6   End station           115597 non-null  object
7   Bike number           115597 non-null  object
8   Member type           115597 non-null  object
dtypes: int64(3), object(6)
memory usage: 7.9+ MB
```



```
df.describe()
```

**Output:**

	Duration	Start station number	End station number
count	115597.000000	115597.000000	115597.000000
mean	1254.649956	31266.213431	31268.042250
std	2914.317998	187.645048	186.194316
min	60.000000	31000.000000	31000.000000
25%	403.000000	31110.000000	31111.000000
50%	665.000000	31213.000000	31214.000000
75%	1120.000000	31301.000000	31238.000000
max	85644.000000	31805.000000	31805.000000

**# Explore and preprocess the dataset**

```
df.isnull().sum()
```

**Output:**

```
Duration          0
Start date        0
End date          0
Start station number  0
Start station      0
End station number  0
End station        0
Bike number        0
Member type        0
dtype: int64
```

```
df=df.drop(['Start date','End date'],axis=1)
```

```
df.head()
```

**Output:**

	Duration	Start station number	Start station	End station number	End station	Bike number	Member type
0	1012	31208	M St & New Jersey Ave SE	31108	4th & M St SW	W00742	Member
1	61	31209	1st & N St SE	31209	1st & N St SE	W00032	Member
2	2690	31600	5th & K St NW	31100	19th St & Pennsylvania Ave NW	W00993	Member
3	1406	31600	5th & K St NW	31602	Park Rd & Holmead Pl NW	W00344	Member
4	1413	31100	19th St & Pennsylvania Ave NW	31201	15th & P St NW	W00883	Member

### # Encoding the categorical values

```
from sklearn.preprocessing import LabelEncoder  
label_encoder=LabelEncoder()  
categorical_columns=['Start station','End station','Bike number','Member type']  
df[categorical_columns]=df[categorical_columns].apply(label_encoder.fit_transform)  
df[categorical_columns].head()
```

### Output:

	Start station	End station	Bike number	Member type
0	85	50	614	1
1	32	33	41	1
2	52	31	836	1
3	52	94	282	1
4	30	21	734	1

### #Split the dataset into features(x)

```
x=df.iloc[:, :-1]
```

x

### Output:

	Duration	Start station number	Start station	End station number	End station	Bike number
0	1012	31208	85	31108	50	614
1	61	31209	32	31209	33	41
2	2690	31600	52	31100	31	836
3	1406	31600	52	31602	94	282
4	1413	31100	30	31201	21	734
...	...	...	...	...	...	...
115592	2179	31110	35	31623	65	716
115593	953	31106	63	31401	18	764
115594	737	31602	93	31401	18	819
115595	514	31111	1	31202	14	946
115596	51962	31111	1	31111	1	636

115597 rows × 6 columns

**#Split the variables into target variables(y)**

```
y=df.iloc[:,-1]
```

y

**Output:**

```
0      1
1      1
2      1
3      1
4      1
..
115592  0
115593  1
115594  1
115595  1
115596  0
Name: Member type, Length: 115597, dtype: int32
```

**# Further split the dataset into training and testing sets**

```
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,random_state=365)
```

**# Feature scaling the data**

```
from sklearn.preprocessing import StandardScaler
```

```
scaler=StandardScaler()
```

```
x_train=scaler.fit_transform(x_train)
```

```
x_test=scaler.transform(x_test)
```

x\_train

**Output:**

```
array([[ 0.06558489,  0.71376786,  0.90403143, -0.26671951,  1.16650355,
        -0.26032801],
       [-0.18886984, -0.28742751, -0.37449522, -0.29892564, -0.35585168,
        -1.12351352],
       [-0.23591478, -0.34600809, -0.94646767, -0.28819026, -0.86330342,
         0.3806874 ],
       ...,
       [-0.09340639, -0.8625823 ,  0.2647681 , -0.29355795, -0.8294733 ,
        -0.59540424],
       [-0.27437488, -0.87855882, -1.11469486, -0.36333788,  1.33565413,
        -0.79572155],
       [ 3.79655747, -0.33535708, -1.1483403 , -0.2720872 ,  1.84310587,
         0.07474823]])
```

x\_test

### Output:

```
array([[ -0.28330311,  1.84810091,  0.76944968,  1.88572304,  1.23416378,
         0.93064948],
       [ 0.07004901,  1.81082236,  1.61058563, -0.27745489, -1.06628411,
         0.37704527],
       [ 1.50406111, -1.40578404, -0.50907698, -1.42613994, -0.49117214,
        -1.27284097],
       ...,
       [-0.37979674, -1.38980752,  1.57694019, -1.40466919,  1.74161552,
        -1.28376737],
       [-0.28021257, -0.22352142, -0.67730417, -0.28819026, -0.86330342,
        -0.31496    ],
       [-0.25033732, -0.2075449 , -1.31656749, -0.33649945, -0.38968179,
         0.90879668]])
```

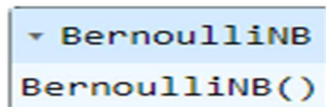
### #Initialise the Bernoulli Naïve Bayes Model

```
from sklearn.naive_bayes import BernoulliNB
model = BernoulliNB ()
```

### #Train the Bernoulli Naïve Bayes Model

```
model.fit(x_train,y_train)
```

### Output:



### #Make predictions

```
y_pred=model.predict(x_test)
y_pred
```

### Output:

```
array([1, 0, 0, ..., 1, 1, 1])
```

### #Evaluate the model

```
from sklearn.metrics import accuracy_score,classification_report
accuracy=accuracy_score(y_test,y_pred)*100
```

```
print("Accuracy of the model is {:.2f}".format(accuracy))
```

**Output:**

```
Accuracy of the model is {:.2f} 82.57093425605537
```

**# Classification Report**

```
from sklearn.metrics import classification_report
class_report=classification_report(y_test,y_pred)
print(f"Classification report: {class_report}")
```

**Output:**

```
Classification report:                precision    recall  f1-score   support

     0      0.58      0.57      0.58      5952
     1      0.89      0.89      0.89     22945
     2      0.00      0.00      0.00         3

 accuracy      0.83      28900
 macro avg      0.49      28900
weighted avg      0.83      28900
```

**#Predict for new data**

```
new_data=[737,31110,63,31623,18,636]
predictions=model.predict([new_data])
print([predictions])
```

**Output:**

```
[array([0])]
```

**# Decoded the output of prediction**

```
decoded_predictions=label_encoder.inverse_transform([predictions])
print(decoded_predictions)
```

**Output:**

```
['Casual']
```

## PROGRAM

### # Import Libraries

```
import pandas as pd
```

```
import numpy as np
```

```
import warnings
```

```
warnings.simplefilter("ignore")
```

### # Load the Dataset

```
df = pd.read_csv("C:\\Users\\Documents\\age.csv")
```

```
df.head()
```

## Output

	name	birthday	title	character_name	character_year	characterage	character_gender	love_interest	release_date	actor_age
0	Ben Platt	24-09-1993	The Politician	Payton Hobart	hs senior	NaN	M	Alice Charles, Astrid Sloan, River Barkley	27-09-2019	26.0
1	Zoe Deutch	10-11-1994	The Politician	Infinity Jackson	hs senior	NaN	F	NaN	27-09-2019	24.0
2	Lucy Boynton	17-01-1994	The Politician	Astrid Sloan	hs senior	NaN	F	Payton Hobart, River Barkley	27-09-2019	25.0
3	Julia Schlaepfer	03-03-1995	The Politician	Alice Charles	hs senior	NaN	F	Payton Hobart, James Sullivan	27-09-2019	24.0
4	Laura Dreyfuss	22-08-1988	The Politician	McAfee Westbrook	hs senior	NaN	F	Skye Leighton	27-09-2019	31.0

```
df.info()
```

## Output

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 243 entries, 0 to 242
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   name                   243 non-null   object
1   birthday               227 non-null   object
2   title                  243 non-null   object
3   character_name         243 non-null   object
4   character_year         111 non-null   object
5   characterage           131 non-null   float64
6   character_gender       243 non-null   object
7   love_interest          132 non-null   object
8   release_date           243 non-null   object
9   actor_age              227 non-null   float64
dtypes: float64(2), object(8)
memory usage: 19.1+ KB
```



```
df.describe()
```

## Output

	characterage	actor_age
count	131.000000	227.000000
mean	16.580153	21.977974
std	1.518877	3.908743
min	10.000000	11.000000
25%	16.000000	20.000000
50%	17.000000	22.000000
75%	17.000000	24.000000
max	25.000000	32.000000

## # Explore and preprocess the data

```
df.isnull().sum()
```

## Output

```
name          0
birthday      16
title         0
character_name 0
character_year 132
characterage   112
character_gender 0
love_interest 111
release_date   0
actor_age      16
dtype: int64
```

## # Preprocessing

### # Extract year from birthday

```
df['birthday'] = pd.to_datetime(df['birthday'])
```

```
df['birthday']
```

Output

```
0      1993-09-24
1      1994-11-10
2      1994-01-17
3      1995-03-03
4      1988-08-22
...
238      NaT
239      1991-10-19
240      1994-11-09
241      1998-04-20
242      NaT
Name: birthday, Length: 243, dtype: datetime64[ns]
```

```
df['birth_year'] = df['birthday'].dt. year
df['birth_year']
```

Output

```
0      1993.0
1      1994.0
2      1994.0
3      1995.0
4      1988.0
...
238      NaN
239      1991.0
240      1994.0
241      1998.0
242      NaN
Name: birth_year, Length: 243, dtype: float64
```

# Drop the null values

```
df = df. drop (['birthday', 'title','character_name', 'character_year', 'characterage',
'love_interest', 'release_date'],
axis=1)
df
```

Output

	name	character_gender	actor_age	birth_year
0	Ben Platt	M	26.0	1993.0
1	Zoe Deutch	F	24.0	1994.0
2	Lucy Boynton	F	25.0	1994.0
3	Julia Schlaepfer	F	24.0	1995.0
4	Laura Dreyfuss	F	31.0	1988.0
...	...	...	...	...
238	Thomas Mitchell Barnet	M	NaN	NaN
239	Kevin Alves	M	28.0	1991.0
240	Asha Bromfield	F	25.0	1994.0
241	Felix Mallard	M	21.0	1998.0
242	Hallee Jones	F	NaN	NaN

243 rows x 4 columns

### # Encoding the categorical values

```
from sklearn.preprocessing import LabelEncoder  
label_encoder = LabelEncoder()  
categorical_columns = ['name', 'character_gender']  
df[categorical_columns] = df[categorical_columns].  
apply(label_encoder.fit_transform)  
df[categorical_columns]
```

### Output

	name	character_gender
0	24	1
1	239	0
2	151	0
3	126	0
4	142	0
...	...	...
238	229	1
239	135	1
240	18	0
241	80	1
242	94	0

243 rows × 2 columns

### # Explore the null values in the data

```
df.isnull().sum()
```

### Output

```
name          0  
character_gender  0  
actor_age     16  
birth_year    16  
dtype: int64
```

## # Replace the null values

```
from sklearn. impute import SimpleImputer

numeric_columns = ['birth_year','actor_age']

df[numeric_columns] = SimpleImputer(strategy='most_frequent').
fit_transform(df[numeric_columns])

df[numeric_columns]
```

## Output

	birth_year	actor_age
0	1993.0	26.0
1	1994.0	24.0
2	1994.0	25.0
3	1995.0	24.0
4	1988.0	31.0
...	...	...
238	1996.0	22.0
239	1991.0	28.0
240	1994.0	25.0
241	1998.0	21.0
242	1996.0	22.0

243 rows × 2 columns

## # Split the data into features (x)

```
x = df[['name', 'character_gender', 'birth_year']]
```

x

## Output

	name	character_gender	birth_year
0	24	1	1993.0
1	239	0	1994.0
2	151	0	1994.0
3	126	0	1995.0
4	142	0	1988.0
...	...	...	...
238	229	1	1996.0
239	135	1	1991.0
240	18	0	1994.0
241	80	1	1998.0
242	94	0	1996.0

243 rows × 3 columns

**# Split the data into target variable (y)**

```
y = df['actor_age']
```

y

**Output**

```
0      26.0
1      24.0
2      25.0
3      24.0
4      31.0
...
238    22.0
239    28.0
240    25.0
241    21.0
242    22.0
Name: actor_age, Length: 243, dtype: float64
```

**# Further split the dataset into training and testing sets**

```
from sklearn. model_selection import train_test_split
```

```
x_train, x_test, y_train, y_test = train_test_split (x, y, test_size=0.2,
random_state=42)
```

**# Intialize the Linear Regression model**

```
from sklearn. linear_model import LinearRegression
```

```
model = LinearRegression ()
```

**# Train the Linear Regression model**

```
model.fit (x_train, y_train)
```

**Output**

```
LinearRegression
LinearRegression()
```

**# Make predictions**

```
y_pred = model. predict(x_test)
```

y\_pred

## Output

```
array([19.75525433, 23.63369918, 20.44610394, 13.18831839, 20.7292717 ,
       19.46859378, 20.24029203, 19.00068209, 26.23148137, 20.32947086,
       23.77648375, 26.83318983, 23.22736137, 19.17661189, 22.06893622,
       20.84317085, 21.83102103, 26.07153311, 19.76921193, 21.12654996,
       21.86640574, 20.73598235, 22.77421559, 20.79357951, 18.45827035,
       20.11660037, 28.28628134, 20.49599376, 21.23635937, 21.0025716 ,
       21.40288227, 20.97014395, 20.53465991, 22.35531006, 20.7046497 ,
       20.68993324, 23.85036094, 20.85675462, 24.00318704, 22.14324788,
       22.91502149, 20.32592862, 20.9635693 , 20.56237497, 20.60576549,
       22.31634543, 19.60239055, 20.65835157, 21.4483462 ])
```

---

## # Evaluate the model

```
from sklearn.metrics import mean_squared_error
mse = mean_squared_error(y_test, y_pred)
print(f"Mean Squared Error: {mse}")
```

## Output

Mean Squared Error: 8.89497933910581

## # Predict the new data

```
new_data = pd.DataFrame({"name": [142], "character_gender":
[1], 'characterage': [17], "b_year": [1996.0]})
predicted_age = model.predict(new_data)
print(f"Predicted Actor Age: {predicted_age [0]}")
```

## Output

Predicted Actor Age: 21.037446484786074